

Pranoy Ghosh

+1 (929) 866-9605 | pranoy.ghosh24@gmail.com | linkedin.com/in/pranoy24 | github.com/pranoyghosh

EDUCATION

New York University – Courant Institute

Sep. 2025 – Expected May 2027

MS in Computer Science — GPA: 3.72/4

New York, USA

- **Coursework:** Machine Learning, Deep Learning, Computer Vision, GPU Programming, Technologies for Finance.

Manipal Institute of Technology

Jul. 2017 – Aug. 2021

B.Tech in Computer and Communication Engineering — CGPA: 9.07/10

Manipal, India

EXPERIENCE

Software Engineer Intern

Jun. 2026 – Aug. 2026

Cohesity Inc.

Santa Clara, USA

- Led containerization of **Throttler**, a replication and archival rate-limiter, while helping migrate core distributed-filesystem services to Kubernetes; implemented its onboarding-framework plugin and **SPIFFE/SPIRE** mutual TLS.
- Built diagnostic tools to verify on-disk chunk/blob metadata integrity and filesystem state; resolved a critical concurrency/locking crash and led root-cause analysis of a customer snapshot/archival failure.

Member of Technical Staff – DevOps / SRE

Aug. 2021 – Jul. 2025

Cohesity Inc.

Bengaluru, India

- Architected a lab-automation platform serving 600+ engineers with 100+ monthly active users, reducing infrastructure provisioning time by 60% through Infrastructure as Code.
- Engineered a Python log parser processing 5M+ diagnostic events daily and a performance analyzer that surfaced I/O bottlenecks in a distributed filesystem, reducing time to diagnosis for field engineers.
- Integrated lab automation with Cohesity's Knowledge Base to generate break-fix scenarios, reducing SRE onboarding time by 50%; served as SME for capacity, garbage collection, and data-healing issues.

PROJECTS

Swin3D-JEPA – Self-Supervised Learning for Physical Simulations

Spring 2026

Prof. Alfredo Canziani and Prof. Yann LeCun, NYU

Python, PyTorch, W&B

- Adapted a hierarchical 3D Swin Transformer to a JEPA objective on 11-channel active-matter simulations; Stage 2+4 features achieved the best linear-probe mean MSE (0.168) across the evaluated self-supervised variants.
- Built automated sweep checks for representation collapse, validation-loss divergence, and incomplete runs; used linear/kNN probes, ablations, and held-out evaluation to rank only stable models, and diagnosed a missing shifted-window attention mask.

Pico LLM – Decoder-Only Transformer

Fall 2025

Prof. Matus Telgarsky, NYU

Python, PyTorch

- Built a causal transformer from scratch with custom attention, positional embeddings, KV caching, attention visualization, and feature-based steering.

Failure-Aware Multi-View Latent Planning

Fall 2026 – Present

Prof Saining Xie, NYU

Python, PyTorch, RLBench

- Designing a reliability-weighted latent planner for robot manipulation under missing, blurred, noisy, stale, or occluded camera views; evaluation isolates reliability weighting using a shared frozen predictor and held-out corruptions.

PUBLICATION

- **Pranoy Ghosh** et al., “Exploring Techniques for Photo-realistic Image Generation from 3D Models: A Deep Learning Approach,” IEEE MysuruCon 2021. [DOI](#)

TECHNICAL SKILLS

Programming: Python, C/C++, SQL, Bash, Go, JavaScript; Git, Linux, debugging.

ML and LLM Systems: PyTorch, TensorFlow, Transformers, Model Context Protocol (MCP), self-supervised learning, linear/kNN probing, ablations, failure analysis, W&B.

Systems and Data: Docker, Kubernetes, AWS, Azure, Terraform, Slurm/HPC, Flask, Dash, Django.